

Whether a Relation Extractor Learns Argument Order Depends on a Corpus Property Nobody Reports

Anonymous ACL submission

Abstract

A directional relation extractor should identify the same entity as subject whichever order the pair is presented in. Nothing in standard practice enforces this, and we show that whether it holds depends on a property of the training corpus that is never reported: the marginal probability that the first-listed argument is the subject. Across three corpora we find that when this prior is skewed (0.708) an unconstrained model takes the positional shortcut and its swap consistency falls to 0.473—below chance; when the prior is near balanced (0.486, 0.544) the model is forced to read content and reaches 0.94–0.97, but never exactness. Holding a corpus fixed and moving only this prior—reordering argument presentation and nothing else—reproduces the effect on two benchmarks, and at a prior of 1.0 swap consistency is exactly 0 while accuracy falls to precisely the test set’s own positional prior. The released state-of-the-art BioREDirect model sits at 0.700 on its own benchmark, where 17.9% of its directed predictions are returned unchanged under exchange—a self-contradiction—and 16.7% are abandoned to the undirected class. We give two mechanisms that make the property exact rather than contingent. Canonicalising argument order before encoding makes the two presentations byte-identical; an antisymmetric readout $s(h_A, h_B) = g(h_A, h_B) - g(h_B, h_A)$, paired with a *symmetric* directedness score, makes it hold at every parameter setting given order-invariant encodings. The first is what makes the property exact—to floating-point precision, whatever the head; the second recovers most of it (0.987) when the encoder stays order-sensitive. Together they cost 0.54% of parameters and 2% of throughput. The guarantee is free rather than beneficial: direction accuracy is statistically indistinguishable from the unconstrained arm on both public benchmarks (0.911 vs 0.912 on BioRED, $p = 0.79$; 0.963 vs 0.966 on SemEval-2010 Task 8, $p = 0.13$). We argue

the right way to obtain a property the evaluation convention already assumes is to construct it, not to hope the corpus happens to supply it, and we release a swap-consistency protocol that needs no new annotation and runs on any existing extractor.

1 Introduction

For directional relations, the label encodes an ordering: *Cause-Effect*($e1, e2$) and *Cause-Effect*($e2, e1$) are distinct classes, and confusing them produces a knowledge-base edge pointing the wrong way. The convention is near-universal, from SemEval-2010 Task 8 (Hendrickx et al., 2010) through TACRED (Zhang et al., 2017) and DocRED (Yao et al., 2019) to recent biomedical work (Lai et al., 2025).

The convention carries an unstated assumption. The same pair can be presented in either order, and the two presentations are semantically identical: if A causes B then presenting (B, A) must yield the mirrored label, so that the entity identified as subject does not change. Nothing in the standard training objective enforces this. Each pair is seen in one order during training and the model is never asked what it would say about the reverse.

We find that the property is therefore not reliably learned—and, more interestingly, that whether it is learned at all depends on a property of the corpus that is never reported. Write π for the marginal probability that the first-listed argument is the subject. Across three corpora with $\pi = 0.708$, 0.544 and 0.486, an unconstrained model trained on order-sensitive input reaches swap consistency 0.473, 0.938 and 0.969 respectively (Figure 1). When π is skewed the model takes the positional shortcut and performs *below chance* on consistency; when π is balanced it is forced to read content and approximates the property without ever attaining it. The released BioREDirect model (Lai et al., 2025) sits at 0.700 on its own benchmark.

This reframes what the architecture is for. The




Figure 1: Swap consistency of an unconstrained model on order-sensitive input, against the training corpus’s positional prior π . The property is acquired only when the corpus withholds the positional shortcut; at $\pi = 0.708$ it falls below chance. Bars span three seeds (biodiversity is a single run). Our canonical arm is at 1.000 for every π , by construction rather than by training.

property is not merely unenforced; it is contingent on a corpus statistic that practitioners neither measure nor control. Constructing it removes the contingency. We give two mechanisms that act at different points. The first removes argument order from the encoder: sorting the two entities canonically before marking them makes the two presentations byte-identical, so consistency holds exactly, for any readout. The second constrains the readout: an antisymmetric head returns $s(h_A, h_B) = -s(h_B, h_A)$ at every parameter setting, which delivers the property whenever the encodings it is given are themselves order-invariant, and recovers most of it (0.987) when they are not. We pair the latter with a *symmetric* directedness score, because whether a relation is directed at all must *not* change under exchange—a distinction a single softmax over {rightward, leftward, undirected} cannot express.

Our claims, each supported on more than one corpus:

1. **Argument-order equivariance is contingent, not guaranteed**, and the contingency is predictable from a measurable corpus statistic (§4.1) that *causes* it: holding a corpus fixed and moving only the statistic reproduces the effect on two benchmarks (§4.2).
2. **Exactness is cheap**. 0.54% additional parameters, 2% of throughput, holding to floating-point precision rather than approximately (§4.4).

3. **It is free**. Direction accuracy is statistically indistinguishable from the unconstrained arm on both public benchmarks (§4.5). We claim no accuracy benefit.

We are explicit about what is not new. Separating relation identity from relation direction dates to Rink and Harabagiu (2010). Antisymmetry by construction is standard in preference learning (Borges et al., 2005; Rigutini et al., 2011) and motivates several knowledge-graph embedding models (Trouillon et al., 2016; Sun et al., 2019). Placing a direction head on a shared encoder is Lai et al. (2025). What we have not found is an argument-order-invariant *input* for relation extraction: the field puts order into the label space, into role markers (Zhou and Chen, 2022; Baldini Soares et al., 2019), or into a constrained permutation set (Woo et al., 2022), rather than removing it from the encoder and recovering it at an equivariant readout.

2 Literature review

Direction in relation extraction. The dominant treatment folds direction into the label set. SemEval-2010 Task 8 (Hendrickx et al., 2010) defines nine relation types in two argument orders plus *Other*; TACRED (Zhang et al., 2017) and DocRED (Yao et al., 2019) order the arguments in the instance instead. The winning SemEval system (Rink and Harabagiu, 2010) already factored the problem, using one classifier for relation type and a second, per type, for direction, so the factorisation itself is not novel. What it lacks, and what we add, is any invariance property: both stages see the arguments in text order and are free to exploit it.

Entity markers. Marking the arguments in the passage is standard and the punctuation we use is borrowed. Zhou and Chen (2022) establish the typed entity marker in which the subject is enclosed in @ and the object in #; Baldini Soares et al. (2019) popularised the idea. We invert the semantics: our markers denote *positions in a canonical ordering*, not roles. A reader familiar with the convention will otherwise assume the standard reading, under which the invariance argument does not hold, so we flag the inversion explicitly.

Symmetry as an architectural constraint. Enforcing $f(a, b) = -f(b, a)$ by construction is old. RankNet (Borges et al., 2005) gives $P(a \succ b) = 1 - P(b \succ a)$ identically; SortNet (Rigutini et al., 2011) enforces the antisymmetry exactly by dual-

neuron weight sharing and proves the constrained network remains a universal approximator—a stronger result than the difference form we use. In knowledge-graph embedding, the inability of symmetric scoring functions to represent antisymmetric relations motivates ComplEx (Trouillon et al., 2016) and RotatE (Sun et al., 2019). Our contribution is not the construction but its application: to a text encoder that must read direction from a passage, paired with a symmetric score for undirectedness, and measured.

Invariance and equivariance. Woo et al. (2022), in scene-graph generation, state the goal in the same vocabulary—equivariant under subject-object ordering while invariant to the permutation—and resolve it in the opposite direction, by constraining the permutation set so the subject always precedes the object. Frame averaging (Puny et al., 2022) is the general umbrella: our canonical sort is a canonicalisation frame giving exact invariance, and our antisymmetric readout is group averaging in the sign representation of S_2 .

Shortcut learning and dataset artefacts. That models exploit superficial statistics rather than the intended signal is well documented (Gururangan et al., 2018; McCoy et al., 2019; Geirhos et al., 2020). Our π statistic is an instance: a positional regularity that is invisible in standard reporting and that determines whether an intended invariance is acquired. Partial-input baselines (Poliak et al., 2018) are the established diagnostic for the analogous problem in inference; our masked-passage control (§4.5) plays that role here.

Consistency probes. Perturbing an input in a way that should have a predictable effect is established methodology (Ribeiro et al., 2020; Elazar et al., 2021). Yuan and Vlachos (2025) show pre-trained encoders are at chance on a symmetry probe built by swapping subject and object in Wikidata triples, and address it with symmetry-aware training objectives; we address it architecturally, so no training signal is required. Woo et al. (2022) introduce a bidirectional diagnostic in vision. Swap consistency (§3.3) is the analogue for directional relation extraction and needs no new annotation.

Biomedical directional extraction. Lai et al. (2025) is our principal baseline. It adds 10,864 manual directionality annotations to BioRED (Luo et al., 2022) and predicts relation type, novelty and direction from three heads on a shared encoder,

with soft prompts and context chunking. Its direction head is a four-way linear classifier over $\{None, Left_to_Right, Right_to_Left, No_Direct\}$ —confirmed here from the released weights—and carries no symmetry constraint.

3 Method

3.1 Task and the property at issue

Given a passage p , an ordered pair of marked mentions (x, y) and a relation type t , predict whether the relation is directed and, if so, which of x, y is its subject. Write $\sigma(p, x, y, t) \in \{x, y, \emptyset\}$. The property is

$$\sigma(p, x, y, t) = \sigma(p, y, x, t), \quad (1)$$

a statement about the identity of the predicted subject, not about the label: the label must mirror precisely so that the identified entity does not change.

3.2 The corpus statistic π

Let π be the fraction of directed training instances whose *first-listed* argument is the subject. A model given order-sensitive input can always achieve accuracy $\max(\pi, 1 - \pi)$ by ignoring the passage. Where π departs from 0.5 this shortcut is rewarded, and the intended content-based solution is correspondingly less necessary. π is trivially computable and, to our knowledge, never reported.

3.3 Swap consistency

We measure Eq. 1 directly: construct the exchanged presentation of each test pair, run the model on both, and record whether the same entity is named subject. The construction is mechanical but must exchange *every* slot in which the input encodes order—role markers and any template stating the pair order.

It is easy to get wrong. Our first attempt on BioREDirect exchanged only the source and target markers, leaving the question order intact, and reported 0.175. That is an artefact: the model reads the textual order of the entities rather than the marker labels, so the manipulation changed the gold answer without changing the argument order of the input. The corrected protocol exchanges both and gives 0.700. We report the error because it is the first thing a careful reader will probe.

3.4 Two routes to exactness

Let h_A, h_B be the contextual representations of the marked spans and c the sentence representation.

Canonical input (invariance). Sort the two entity strings by a fixed total order—case-folded lexicographic—and mark the first @ · @ and the second # · # in place. The token sequence is then a function of the *unordered* pair: the two presentations are byte-identical and Eq. 1 holds for any readout. Decoding to a label happens outside the model, by comparing the @-marked entity to the first argument as given. Canonicalisation also disarms π . Exactness does not depend on the canonical prior being balanced—the two presentations are the same bytes whatever it is—but the shortcut also stops paying, because which entity occupies the first slot is now decided by the alphabet rather than by the corpus. Measured, the prior moves to 0.484 on SemEval (from 0.544) and 0.502 on BioRED (from 0.486): alphabetical order carries essentially no information about which argument is the subject, which is what makes it a usable canonical form.

Antisymmetric readout (equivariance). Independently of the input,

$$s_{\text{dir}} = g([h_A; h_B; c; e_t]) - g([h_B; h_A; c; e_t]), \quad (2)$$

$$s_{\text{und}} = h([h_A; h_B; c; e_t]) + h([h_B; h_A; c; e_t]), \quad (3)$$

with e_t a small embedding of the relation type and g, h two-layer networks. Eq. 2 is exactly antisymmetric and Eq. 3 exactly symmetric, at every parameter setting, before and during training.

The pairing carries the argument. *Which* entity is the subject must invert under exchange; *whether the relation is directed at all* must not. A single softmax over {rightward, leftward, undirected} expresses neither constraint: exchanging the arguments can move mass between rightward and leftward *and* into or out of undirected, so the model may hold two incompatible beliefs about one pair. We verify both properties numerically before training; residuals are 0 to machine precision.

Keeping the relation out of the encoder. The relation type enters only at the head, through e_t ; the encoder never sees the relation string. This makes the invariance argument airtight—direction cannot re-enter through the relation surface form—and lets one small head serve every relation in the schema.

3.5 Model, cost and protocol

The encoder is BiomedBERT-base (Gu et al., 2021) in all arms. The direction head is 594,465 parameters, 0.54% of the encoder, costing 2% of

Corpus	π	Uncon. swap	Ours
Biodiversity	0.708	0.473	1.000
SemEval-2010	0.544	0.938	1.000
BioRED	0.486	0.969	1.000
BioREDirect (released)		0.700	—

Table 1: Swap consistency of the order-sensitive unconstrained arm against the training-corpus prior π , with our canonical arm for reference.

CPU throughput (33.0 against 33.8 pairs/s on eight threads). Arms within a comparison share encoder, optimiser, schedule, epochs and data, differing only on the axis under test. We report means over three seeds.

3.6 Corpora

BioRED (Luo et al., 2022) with the directionality annotations of Lai et al. (2025): 1,163 test pairs carrying a relation, 1,045 directed (527/518) and 118 undirected; $\pi = 0.486$. Clean mentions and passages are recovered from the PubTator source, using the smallest sentence window containing both entities as the baseline’s own chunking does.

SemEval-2010 Task 8 (Hendrickx et al., 2010): 7,991 training and 2,713 test instances after removing pairs with identical mention strings; 2,260 directed test instances and 453 *Other*, treated as undirected. General domain. $\pi = 0.544$.

Biodiversity verification set: 437 expert-graded candidate triples from a deployed literature-mining pipeline, plus 84 human-annotated direction items over 79 relation surface forms. $\pi = 0.708$ —the skewed case. Small, and treated accordingly.

4 Results

4.1 Equivariance is contingent on π

Table 1 and Figure 1 are the central result. An unconstrained model on order-sensitive input—standard practice—acquires the equivariance only when the corpus forces it to. At $\pi = 0.708$ the positional shortcut is available and consistency falls to 0.473, below chance. At $\pi \approx 0.5$ the shortcut is worthless and consistency reaches 0.94–0.97, but never 1.

We derived this relationship from the biodiversity and BioRED corpora and used it to predict SemEval before running it: at $\pi = 0.544$, mildly skewed, consistency should fall between and nearer the balanced end. It does.

π	Dir. acc.	Swap	BioRED swap
0.50	0.9608	0.9889	0.9842
0.60	0.9604	0.9805	—
0.70	0.9582	0.9628	0.9732
0.80	0.9502	0.9449	—
0.90	0.9272	0.8779	0.8789
1.00	0.5845	0.0000	—
<i>canonical + symmetric, same training sets</i>			
0.50	0.9646	1.0000	—
1.00	0.9650	1.0000	—

Table 2: π imposed on the training set by reordering argument presentation and nothing else. Two seeds per point; standard deviations are at most 0.017, and both seeds of every point are in Appendix B.

4.2 π causes it: a controlled intervention

Three corpora with a mechanism attached is still an observation. To make the claim causal we hold the corpus fixed and move π directly. For a target π we rewrite the SemEval training set by changing *only the order in which the two arguments are presented*, flipping a chosen subset of pairs. The passages, the pairs, the relation types, the gold subjects and the test set are all untouched—the generator asserts that the pair multiset, the passages and the directedness flags are unchanged before it writes—so π is the single manipulated variable. Two seeds per point.

Table 2 and Figure 2 give the result. Swap consistency falls monotonically with π , from 0.989 at 0.50 to 0.878 at 0.90, while direction accuracy barely moves. That gap is the practical warning in one line: between $\pi = 0.5$ and $\pi = 0.9$ a model loses eleven points of consistency for three points of accuracy, so a leaderboard reporting accuracy would show almost nothing wrong.

The endpoint is unambiguous. At $\pi = 1.00$, where every training pair is presented subject-first, swap consistency is exactly 0.0000: the model names the first-listed entity every time, so exchanging the arguments always changes its answer. Its accuracy is 0.5845, which is exactly the test set’s own positional prior to four decimal places—the signature of a model that has learned the shortcut and nothing else. Trained on the identical data, the canonical arm is untouched: 0.9650 accuracy and 1.000 consistency at $\pi = 1.00$ against 0.9646 and 1.000 at $\pi = 0.50$. Immunity to the corpus’s positional prior is what the construction buys.

Repeating the manipulation on BioRED reproduces the curve (0.984 $\rightarrow$ 0.973 $\rightarrow$ 0.879 at $\pi = 0.50, 0.70, 0.90$), so the effect is not a property of one dataset.




Figure 2: The same corpus at six priors. Consistency degrades while accuracy holds, until the shortcut becomes worth more than reading the passage. The constrained arm, trained on the same data, is flat at 1.000 at both extremes.

π sets the price, not the outcome. One comparison refines the account. SemEval forced to $\pi = 0.70$ keeps consistency at 0.963, whereas the biodiversity corpus at its natural $\pi = 0.708$ collapses to 0.473. π therefore fixes what the shortcut is *worth*— $\max(\pi, 1 - \pi)$ on the training distribution—but whether the model takes it depends on what reading the passage is worth instead. In SemEval a content solution is worth 0.966 against a 0.70 shortcut and wins easily. In the biodiversity data the best content solution we obtain is 0.869, and the relation surface form adds almost nothing on top of position: a lookup table from the relation form predicts the subject at 0.7152 against 0.7075 for answering “first-listed” unconditionally. The shortcut is nearly competitive with everything else available, and it wins. The sweep shows the same crossover within a single corpus: SemEval degrades gently while the shortcut is worth less than 0.966, and collapses once it is worth more.

4.3 What the failure looks like

On the released BioREDirect model, of the 854 BioRED test pairs for which it predicts a direction, 65.3% are correctly mirrored under exchange, 17.9% are returned unchanged—naming a different entity as subject depending on presentation order—and 16.7% are abandoned to the undirected class.

The failure is specific. Relation detection barely moves under exchange (F_1 0.7926 $\rightarrow$ 0.7861) while relation-plus-direction falls 0.5550 $\rightarrow$ 0.4013. What the model knows about *whether* two entities are related survives; what it knows about

which way does not.

4.4 The guarantee holds exactly

Across every arm using canonical input, swap consistency is 1.000 and the maximum deviation between a prediction and its mirror is 2.98×10^{-8} —floating-point precision, not approximation. This holds whichever head sits on top, because the encoder receives identical bytes in both presentations.

The antisymmetric readout does not by itself deliver the same thing, and it is worth being exact about why. Eq. 2 is antisymmetric in its two arguments at every parameter setting, so it guarantees the property *conditional on* h_A and h_B being unchanged by the exchange. An order-sensitive encoder breaks that condition: it genuinely sees different tokens when the markers move, so h_A and h_B are not the same vectors. Empirically the head recovers most of the gap but never closes it— 0.987 ± 0.002 on SemEval against 0.938 ± 0.028 for the unconstrained head, and 0.969 on the biodiversity set against 0.473. The two mechanisms are therefore not interchangeable: canonicalisation is what makes the property exact, and the antisymmetric head is what makes it nearly free to approximate when canonicalisation is unavailable—for instance where argument order carries information the task needs.

4.5 Accuracy: the guarantee is free

Table 3 gives the head-to-head. On both public benchmarks the canonical-symmetric arm and the order-sensitive unconstrained arm are indistinguishable. Summarising each arm by the majority vote of its three seeds and comparing per item, the unconstrained arm is nominally ahead on both—by 0.002 on BioRED (6 vs 8 discordant items, McNemar with continuity correction $p = 0.79$) and by 0.004 on SemEval (13 vs 23 discordant, $p = 0.13$)—and neither gap approaches significance. **We claim no accuracy benefit from the construction;** we claim it costs nothing.

The 0.766 of the released BioREDirect model is not directly comparable to either arm: it performs a joint task (relation, novelty, direction) with soft prompts and multi-chunk pooling, so the gap reflects system configuration rather than the design axis under test. We report it for context, not as evidence for our design.

Two controls support that the models read the passage rather than exploiting surface regularities. Withholding the relation type entirely changes

Corpus	Arm	Dir. acc.	Swap
BioRED	BioREDirect (released)	0.766	0.700
	text + unconstrained	0.912	0.969
	canonical + sym.	0.911	1.000
SemEval	text + unconstrained	0.966	0.938
	canonical + sym.	0.963	1.000

Table 3: Direction accuracy on gold-directed pairs and swap consistency. All BioRED and SemEval figures are means over three seeds; per-seed values are in Appendices A and B.

BioRED accuracy from 0.909 to 0.910, so the head does not lean on that feature. Masking the passage and presenting only the two entity names drops the biodiversity direction model to 0.464—chance—so it is not recalling which entities typically stand in which role.

4.6 Decomposing the two mechanisms

Table 4 isolates each design choice, on the skewed biodiversity set where the distinction bites and on SemEval, where three seeds per cell and a 453-instance undirected class make the comparison stable.

The design separates into two main effects with no interaction. **Consistency is decided by the input format:** both canonical cells are exactly 1.000, and neither text cell reaches it whatever the head. **Undirectedness is decided by the head:** on SemEval the symmetric branch gains ≈ 0.043 AUC under both input formats. **Accuracy is decided by neither:** the four SemEval cells span 0.0036, within about two standard deviations of seed noise, and on biodiversity no difference reaches significance at $n = 84$ (canonical input $p = 0.239$; symmetric head $p = 1.000$).

One result does not replicate. On the 84-item set the unconstrained head looks like the better undirectedness detector (0.811 vs 0.748); on SemEval that reverses cleanly and consistently. The biodiversity set has too few undirected items to support the comparison, and we report the disagreement rather than the convenient half of it.

4.7 The applied setting

On the biodiversity verification set the same model raises F_1 from 0.816 to 0.902 against a deployed sentence-level filter’s own recorded decisions (McNemar $p = 2.3 \times 10^{-4}$, 437 items) at unchanged inference cost. The honest characterisation is robustness to prevalence shift rather than raw accu-

Input	Head	Acc.	Und. AUC	Swap
<i>Biodiversity</i> , $n = 84$, $\pi = 0.708$, one run				
canon.	sym.	0.869	0.748	1.000
canon.	uncon.	0.857	0.787	1.000
text	sym.	0.810	0.718	0.969
text	uncon.	0.786	0.811	0.473
<i>SemEval</i> , $n = 2,260$, $\pi = 0.544$, 3 seeds				
canon.	sym.	0.9625	0.947	1.0000
canon.	uncon.	0.9634	0.904	1.0000
text	sym.	0.9661	0.947	0.9869
text	uncon.	0.9658	0.907	0.9383

Table 4: Design ablation on both corpora. Canonical input fixes swap consistency at floating-point exactness; the symmetric head fixes undirectedness detection; neither moves accuracy. SemEval figures are means over three seeds; per-seed values are in Appendix A.

racy: across blocks spanning prevalence 0.31 to 0.85, one fixed threshold gives an F_1 spread of 0.070 against the filter’s 0.169.

5 Discussion and conclusion

The measurement is the part we would most like reused. It costs nothing—no annotation, no retraining, one extra forward pass over an existing test set—and on the evidence here it is worth running on any released directional extractor.

The finding it produces is not that published systems are careless. It is that a property the evaluation convention silently assumes is delivered only when the corpus happens to make the alternative unprofitable. π is trivial to compute, is not reported in any dataset paper we examined, and predicts whether the intended invariance will be acquired. A corpus with a skewed positional prior will produce models that look accurate and are incoherent under exchange; a balanced one will not, and neither the model card nor the leaderboard will distinguish them.

The architectural response is modest in isolation and we have said so. The factorisation is Rink and Harabagiu’s; the antisymmetric construction is Rigutini et al.’s; the markers are Zhou and Chen’s with the semantics inverted. What is new is the combination, the measurement that motivates it, and the observation that the property is otherwise left to chance.

Two findings cut against a clean story and belong in the record. The first is a disagreement between our two ablations. On the 84-item biodiversity set the unconstrained head detects undirectedness *better* than ours, which we first read as the symmetry

constraint costing flexibility in the undirected class; on SemEval, with 453 undirected instances and three seeds per cell, the effect reverses cleanly and the symmetric branch wins by 0.043 AUC under both input formats. We take the larger measurement to be the reliable one and report the smaller because it is what a $n = 84$ ablation can do to a conclusion. The second is that no accuracy difference between design cells is significant on either corpus—only the consistency differences are structural, and that is the whole claim.

Limitations. The human direction gold for the applied setting is 84 items: enough to show the head reads the passage, not enough to separate it from a good dependency-parse rule, which we estimate would need on the order of 2,700 annotations. BioREDirect performs a joint task while our arms perform direction given the pair; we control for the most obvious confound (§4.5) but not all, and we therefore do not treat that comparison as evidence for the design. Our context window is the infix span where the baseline uses prefix, suffix and infix with max-pooling. The π relationship is established by controlled manipulation on two corpora and observed on a third; what the intervention does not settle is the functional form, since six points on one corpus and three on another constrain the shape only loosely, nor how π interacts with the many other ways a corpus can differ. Finally, the curated relation vocabulary shows that many relation surface forms carry no direction at all—12 of 79 annotated forms were marked undirected by the curator—so direction is defined only on a subset of extracted edges, and a system reporting it everywhere is overclaiming.

References

- Livio Baldini Soares, Nicholas FitzGerald, Jeffrey Ling, and Tom Kwiatkowski. 2019. [Matching the blanks: Distributional similarity for relation learning](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2895–2905, Florence, Italy. Association for Computational Linguistics.
- Chris Burges, Tal Shaked, Erin Renshaw, Ari Lazier, Matt Deeds, Nicole Hamilton, and Greg Hullender. 2005. Learning to rank using gradient descent. In *Proceedings of ICML*, pages 89–96.
- Yanai Elazar, Nora Kassner, Shauli Ravfogel, Abhishava Ravichander, Eduard Hovy, Hinrich Schütze, and Yoav Goldberg. 2021. Measuring and improving

587	consistency in pretrained language models. <i>Transactions of the Association for Computational Linguistics</i> , 9:1012–1031.	642
588		643
589		644
590	Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A Wichmann. 2020. Shortcut learning in deep neural networks. <i>Nature Machine Intelligence</i> , 2(11):665–673.	645
591		646
592		
593		647
594		648
595		649
596	Yu Gu, Robert Tinn, Hao Cheng, Michael Lucas, Naoto Usuyama, Xiaodong Liu, Tristan Naumann, Jianfeng Gao, and Hoifung Poon. 2021. Domain-specific language model pretraining for biomedical natural language processing. <i>ACM Transactions on Computing for Healthcare</i> , 3(1):1–23.	650
597		
598		651
599		652
600		653
601		654
602	Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel Bowman, and Noah A Smith. 2018. Annotation artifacts in natural language inference data. In <i>Proceedings of NAACL-HLT</i> , pages 107–112.	655
603		656
604		657
605		658
606		659
607	Iris Hendrickx, Su Nam Kim, Zornitsa Kozareva, Preslav Nakov, Diarmuid Ó Séaghdha, Sebastian Padó, Marco Pennacchiotti, Lorenza Romano, and Stan Szpakowicz. 2010. SemEval-2010 task 8: Multi-way classification of semantic relations between pairs of nominals. In <i>Proceedings of the 5th International Workshop on Semantic Evaluation</i> , pages 33–38.	660
608		661
609		662
610		663
611		664
612		
613	Po-Ting Lai, Chih-Hsuan Wei, Shubo Tian, Robert Leaman, and Zhiyong Lu. 2025. Enhancing biomedical relation extraction with directionality. <i>Bioinformatics</i> , 41(Supplement 1):i68–i76.	665
614		666
615		667
616		668
617	Ling Luo, Po-Ting Lai, Chih-Hsuan Wei, Cecilia N Arighi, and Zhiyong Lu. 2022. BioRED: a rich biomedical relation extraction dataset. <i>Briefings in Bioinformatics</i> , 23(5):bbac282.	669
618		670
619		671
620		672
621	Tom McCoy, Ellie Pavlick, and Tal Linzen. 2019. Right for the wrong reasons: Diagnosing syntactic heuristics in natural language inference. In <i>Proceedings of ACL</i> , pages 3428–3448.	673
622		674
623		675
624		
625	Adam Poliak, Jason Naradowsky, Aparajita Haldar, Rachel Rudinger, and Benjamin Van Durme. 2018. Hypothesis only baselines in natural language inference. In <i>Proceedings of *SEM</i> , pages 180–191.	676
626		677
627		678
628		
629	Omri Puny, Matan Atzmon, Heli Ben-Hamu, Ishan Misra, Aditya Grover, Edward J Smith, and Yaron Lipman. 2022. Frame averaging for invariant and equivariant network design. In <i>International Conference on Learning Representations (ICLR)</i> .	679
630		
631		680
632		681
633		682
634	Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. Beyond accuracy: Behavioral testing of NLP models with CheckList. In <i>Proceedings of ACL</i> , pages 4902–4912.	683
635		684
636		685
637		686
638	Leonardo Rigutini, Tiziano Papini, Marco Maggini, and Franco Scarselli. 2011. SortNet: Learning to rank by a neural preference function. <i>IEEE Transactions on Neural Networks</i> , 22(9):1368–1380.	687
639		688
640		689
641		690
		691
		692
	Bryan Rink and Sanda Harabagiu. 2010. UTD: Classifying semantic relations by combining lexical and semantic resources. In <i>Proceedings of the 5th International Workshop on Semantic Evaluation</i> , pages 256–259.	
	Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. 2019. RotatE: Knowledge graph embedding by relational rotation in complex space. In <i>International Conference on Learning Representations (ICLR)</i> .	
	Théo Trouillon, Johannes Welbl, Sebastian Riedel, Éric Gaussier, and Guillaume Bouchard. 2016. Complex embeddings for simple link prediction. In <i>Proceedings of ICML</i> , pages 2071–2080.	
	Sangmin Woo, Junhyug Noh, and Kangil Kim. 2022. Tackling the challenges in scene graph generation with local-to-global interactions. <i>IEEE Transactions on Neural Networks and Learning Systems</i> . ArXiv:2106.08543.	
	Yuan Yao, Deming Ye, Peng Li, Xu Han, Yankai Lin, Zhenghao Liu, Zhiyuan Liu, Lixin Huang, Jie Zhou, and Maosong Sun. 2019. DocRED: A large-scale document-level relation extraction dataset. In <i>Proceedings of ACL</i> , pages 764–777.	
	Zhangdie Yuan and Andreas Vlachos. 2025. Capturing symmetry and antisymmetry in language models through symmetry-aware training objectives. <i>arXiv preprint arXiv:2504.16312</i> .	
	Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli, and Christopher D. Manning. 2017. Position-aware attention and supervised data improve slot filling . In <i>Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing</i> , pages 35–45, Copenhagen, Denmark. Association for Computational Linguistics.	
	Wenxuan Zhou and Muhao Chen. 2022. An improved baseline for sentence-level relation extraction. In <i>Proceedings of AACL-IJCNLP</i> . ArXiv:2102.01373.	

A Per-seed values

Table 5 gives every seed behind the SemEval means in Table 4. Swap consistency for the two canonical cells is exactly 1.0000 in every run; the largest deviation between any prediction and its mirror is 2.98×10^{-8} across all six, which is the resolution of the float comparison rather than a measured discrepancy. The undirectedness AUC carries the largest seed variance of any quantity we report ($\sigma \approx 0.02$ – 0.03), which is why the head effect is stated as a difference of means over three seeds rather than from a single run: on seed 1 alone it would have read as 0.98 against 0.93, roughly twice its stable size.

Input	Head		s1	s2	s3
canon.	sym.	acc	0.9650	0.9615	0.9611
		swap	1.0000	1.0000	1.0000
		AUC	0.9823	0.9245	0.9354
canon.	uncon.	acc	0.9642	0.9624	0.9637
		swap	1.0000	1.0000	1.0000
		AUC	0.9297	0.8865	0.8951
text	sym.	acc	0.9668	0.9681	0.9633
		swap	0.9894	0.9863	0.9850
		AUC	0.9794	0.9238	0.9365
text	uncon.	acc	0.9642	0.9673	0.9659
		swap	0.9588	0.9066	0.9496
		AUC	0.9298	0.8952	0.8962

Table 5: SemEval-2010 Task 8, all three seeds per cell.

Arm		s1	s2	s3
canon. + sym.	acc	0.9091	0.9100	0.9129
	swap	1.0000	1.0000	1.0000
text + uncon.	acc	0.9120	0.9129	0.9100
	swap	0.9809	0.9694	0.9569

Table 6: BioRED test, 1,045 gold-directed pairs, all three seeds.

B Per-seed values for the π intervention

Table 7 gives both seeds of every point in Table 2. The $\pi = 1.00$ row is worth noting: the two seeds agree to four decimals on both metrics, which is what a degenerate solution looks like—there is only one way to answer “first-listed” every time, so independent initialisations land on identical behaviour rather than merely similar behaviour.

Corpus	π	accuracy		swap	
		s1	s2	s1	s2
SemEval	0.50	0.9611	0.9606	0.9925	0.9854
	0.60	0.9593	0.9615	0.9832	0.9779
	0.70	0.9580	0.9584	0.9597	0.9659
	0.80	0.9478	0.9527	0.9420	0.9478
	0.90	0.9226	0.9319	0.8659	0.8898
	1.00	0.5845	0.5845	0.0000	0.0000
BioRED	0.50	0.9053	0.9062	0.9866	0.9818
	0.70	0.9072	0.9072	0.9780	0.9684
	0.90	0.8842	0.8813	0.8746	0.8833

Table 7: The controlled π sweep, both seeds.