

Which Pair Is It About?

Pair-Conditioned Verification of Literature-Mined Biotic Interactions

Anonymous ACL submission

Abstract

Literature-mining pipelines that populate biotic-interaction databases propose candidate (taxon, relation, taxon) triples by co-occurrence, then decide which to keep with a *sentence-level* classifier: does this passage describe an interaction? A candidate is retrieved precisely because its passage names two taxa near an interaction term, so that question is answered before it is asked; what needs deciding is whether the passage asserts an interaction between *these two*. We test the claim with a controlled comparison: one encoder, one 48k-row teacher-labelled corpus, three seeds, differing only in what the model is shown. On 437 expert-graded candidates, telling the model which pair the candidate is about raises AUPRC from 0.851 to 0.919 and F1 from 0.775 to 0.871 (75 items fixed against 22 broken, McNemar $p = 1.3 \times 10^{-7}$), at every one of 99 thresholds; adding the relation term to the query buys nothing further. The same comparison on the public BioRED corpus, with candidates built the same way, gives $0.783 \rightarrow 0.876$ ($p = 4.5 \times 10^{-71}$), and the gain lies where the argument puts it: none on sentences that name two entities, +0.106 on sentences that name more. We analyse the errors that remain, test deterministic candidate rules, and report the deployed model: a 110M-parameter pair-conditioned verifier that also labels interaction direction in four categories and runs at 33 candidates per second on a CPU.

1 Introduction

Biotic-interaction databases such as GloBI (Poe-len et al., 2014) are populated in part by text mining. A retrieval system proposes candidate triples (s_1, r, s_2) together with the passage that allegedly supports them, and something must decide which candidates are worth a curator’s attention. The natural choice, and the common one, is a sentence-level interaction classifier: show it the passage and ask whether it describes a biotic interaction.

That question is answered before it is asked. Candidates are proposed by co-occurrence: a passage naming two taxa near an interaction term yields a candidate whether or not it asserts that interaction between those two. Every candidate therefore arrives already satisfying, lexically, the property a sentence-level classifier was trained to detect — in the pipeline we study, on all 175,588 candidates of its retrieval pool (§2). What actually needs deciding is narrower: whether the passage asserts an interaction between *these two* taxa. When a passage names three organisms it can support one candidate and refute the other two, and a classifier that sees only the passage must give all three the same answer.

We make this argument testable rather than rhetorical. We train the same 110M-parameter encoder on the same teacher-labelled corpus three times — reading the passage alone, the passage with the candidate pair, and the passage with the whole candidate triple — and compare them on a 437-item benchmark graded by a domain expert (§4). Everything except the input is held fixed, so differences between them are attributable to the formulation and to nothing else. Telling the model which pair it is judging wins by a wide and significant margin; telling it the relation term as well does not add to it. We then rebuild the setting on BioRED (Luo et al., 2022), a public corpus with gold entity and relation annotations, and the result replicates there with the gain in the predicted place: absent on sentences that name only two entities, large on sentences that name more.

Contributions.

- A controlled comparison of sentence-level, pair-conditioned and triple-conditioned verification on one encoder and one corpus, replicated on a public benchmark, with the gain located where the argument says it should be (§4).

- An account of what the pair query buys — disambiguation between the pairs a passage names, and a check on candidates whose arguments are not organisms — and a bound showing that checking a triple’s components separately cannot reach the largest error class (§4, §5).
- Deterministic candidate rules, validated on 42k training rows before evaluation, measured as a small precision gain rather than a large one (§6).
- A deployed configuration with a fixed threshold, four-way direction output and CPU throughput, and five negative results (§7, Appendix B).

2 Setting

Where candidates come from. BiotXplorer (Ruch et al., 2024) proposes candidate biotic-interaction triples from the literature indexed by SIB Literature Services (Gobeill et al., 2020). For each passage it matches taxon surface forms against a taxonomy and interaction surface forms against a vocabulary derived from the Relation Ontology, and emits every resulting (s_1, r, s_2) combination with the passage, the matched surface strings, and their canonical forms. This is the application we build for, and we use it only as the source of candidates: no component of it is a baseline in this paper.

The problem this paper addresses was diagnosed by BiotXplorer’s own authors. Ruch et al. (2024) graded 100 random BiotXplorer triples, found that the system identified the interacting species with a precision of 31%, and named the main cause of error: “instances where passages listed multiple species, which can be automatically filtered out.” We take the diagnosis as given and question only the remedy. On our benchmark, passages naming three or more taxa hold 154 of the 246 interacting pairs (63%), so filtering them out would discard most of what the pipeline exists to find. The alternative we test is to keep those passages and decide each candidate in them separately.

Why the sentence-level question is constant here. We pulled the retrieval pool in full: 175,588 candidate rows over 165,325 distinct passages and 115,404 distinct triplet keys. Both taxon surface forms and the interaction surface form are non-empty on all 175,588 rows, because the presence of an interaction term is the pool’s construction

filter, and an independent lexical scan for GloBI interaction vocabulary fires on 99.82% of the same passages. A classifier trained to separate passages with interaction language from passages without it is, on this distribution, being asked about a constant. The missing input is not an annotator artefact but the query the pipeline itself issued.

Benchmark. One benchmark of 437 rows, 246 positive, in three blocks of different provenance; we report every block rather than pooling them into a number that hides their disagreement. BIOTX100 is retrieved candidates, each the top-ranked candidate for its passage, graded by a domain expert on four axes: whether each taxon surface form was correctly resolved, whether the relation was correct, and whether the whole triple is supported by its passage. Half were drawn from pairs already recorded in GloBI and half at random, and after removing three training duplicates it contributes 97 rows: 50 GloBI-matched with 43 positive, 47 random with 32 positive. It is therefore enriched for known interactors and for top-ranked candidates, which is why its positive rate is far above the 31% Ruch et al. (2024) measured on random triples at any rank. REJECT50 is candidates the retrieval pipeline discarded, curated for the ones it should have kept: 41 clean rows, 9 positive. It is the paper’s recall-side instrument. TEST299 is 299 pair-level items from expert curation, 0.542 positive, and is where passages most often name several taxa.

The blocks do not carry the same gold. BIOTX100 is the only block graded at the level of the whole triple; TEST299 was curated pair-level and its stored relation term was attached by retrieval rather than by an annotator. We therefore evaluate everywhere at the level all three share — *do these two taxa interact in this passage* — and use the triple-level grades only to diagnose (§5). This understates what triple conditioning buys, since a candidate with the right pair and the wrong relation counts as correct, and we prefer understating it to scoring two blocks against a relation term nobody checked.

Query representation. Each candidate carries a canonical string and the surface form matched in the passage. On BIOTX100 these differ for s_1 , s_2 and r on 34, 38 and 59 of 100 rows, and the canonical string is absent from the passage on 31, 27 and 54. We always query with the surface form, which is what the model can see.

Evaluation protocol. Thresholds are where this kind of evaluation goes wrong most easily, so the policy is fixed in advance. A model’s reported operating point is either **pre-specified** — chosen on its own training-development split, never on the benchmark — or **block-held-out**: fitted on two of the three source blocks and applied to the third, so that every row is scored by a threshold that never saw it and every row is used. Comparisons between models use the same policy for both. Intervals are 20,000-sample bootstraps or Clopper–Pearson; McNemar tests (McNemar, 1947) carry a continuity correction. Where three seeds exist we report the per-seed mean and standard deviation, and where we report the average of three checkpoints’ probabilities we call it an ensemble. The benchmark ships an `in_train` flag on 9 rows; a 5-gram Jaccard scan against all 48,338 training passages finds 3 more at Jaccard 1.000, all in BIOTX100. All 12 are excluded throughout.

The gold is pair-level and therefore order-free: *do these two taxa interact*, in either direction. A model that reads the candidate in order is scored the same way, as the higher of its two scores over the two argument orders — the probability that the pair interacts in at least one direction. We fix this before scoring because the alternative penalises a model for rejecting a correctly paired candidate whose stated relation runs backwards, which the gold does not count as an error. The sentence-level model reads no candidate and is order-free already.

Two things about this benchmark’s history should be on the record. It is not untouched: during the development of the verifier it was scored for roughly a hundred model configurations, so the comparisons in this paper are between recipes chosen with it in view, and the controlled comparison of §4 is the result we would expect to survive that. And one gold label has been changed since the benchmark was frozen: a TEST299 candidate pairing *Acanthamoeba* with *Pseudomonas*, revised from positive to negative after the curator, annotating direction without model scores, recorded that the two co-occur in a shared host rather than interacting.

3 Two formulations, one corpus

The teacher. Labels come from a locally hosted 32B open-weight instruction model (Qwen Team, 2025), prompted zero-shot at temperature 0 and asked, per candidate, whether the passage supports

this triple. Its labels therefore vary across candidates that share a passage. Against 150 human-graded items its mislabel rate is $22/150 = 14.7\%$, Clopper–Pearson [9.4%, 21.4%], split 10 false positives and 12 false negatives.

The corpus. 48,338 teacher-labelled rows, 0.343 positive, drawn from the same retrieval pool as the benchmark (disjoint passages): a 20,000-row base sample, 16,546 *pair-contrast* rows, 6,201 actively sampled rows, and 5,591 rows with the candidate’s argument order reversed. The pair-contrast block is what gives the corpus its point: 2,017 passages enter as several rows that are identical in the passage and differ only in the candidate and the label.

Three models, one recipe. All three are cross-encoders on BiomedBERT-base (Gu et al., 2021) (microsoft/BiomedNLP-BiomedBERT-base-uncased-110M parameters), distilled from the teacher’s labels (Buciluă et al., 2006; Hinton et al., 2015) with the same recipe: AdamW at learning rate 2×10^{-5} and weight decay 0.01, a linear schedule with 10% warm-up, batch size 32, inputs of at most 256 wordpieces, three epochs, gradient clipping at 1.0, and the checkpoint with the best development AUPRC, on a 10% development split grouped by taxon pair so that no pair spans training and development. Three seeds each; one run takes about four minutes on a single A100. They differ only in their input. The **sentence-level baseline** reads the passage alone. It is the formulation a pipeline reaches for first, and on rows that share a passage it is shown the same input with different labels, which it cannot fit. The **pair-conditioned** model reads the two candidate taxa, s_1 [SEP] s_2 , as segment A and the passage as segment B (Nogueira and Cho, 2019). The **triple-conditioned verifier** reads the whole candidate, s_1 [SEP] r [SEP] s_2 . The middle model separates what the pair buys from what the relation adds.

4 Results

4.1 Conditioning on the pair

Table 1 gives the comparison. Telling the model which pair a candidate is about raises per-seed AUPRC from 0.851 ± 0.006 to 0.919 ± 0.005 and, at block-held-out thresholds, F1 from 0.775 to 0.871, precision and recall both. On decision correctness the pair-conditioned model fixes 75 items and breaks 22, McNemar $p = 1.3 \times 10^{-7}$. The advantage does not depend on where we cut: it

System	AUPRC	P	R	F1
Accept everything	0.563	0.563	1.000	0.720
Sentence-level	0.851 $\pm$.006	0.730	0.825	0.775
Pair-conditioned	0.919$\pm$.005	0.849	0.894	0.871
Triple-conditioned	0.909 $\pm$.004	0.820	0.886	0.852

Table 1: The controlled comparison on 437 rows: one encoder, one corpus, only the input differs. AUPRC is the per-seed mean $\pm$ standard deviation over three seeds; P/R/F1 are three-checkpoint ensembles at block-held-out thresholds. Against the sentence-level model, the pair-conditioned one fixes 75 items and breaks 22 (McNemar $p = 1.3 \times 10^{-7}$) and the triple-conditioned one fixes 74 and breaks 32 ($p = 6.8 \times 10^{-5}$); triple against pair, 20 fixed and 31 broken ($p = 0.16$). Per-seed values are in Appendix A.

Candidates	n	Sentence	Pair	Triple
passage names ≤ 2 taxa	161	0.876	0.918	0.879
passage names ≥ 3 taxa	276	0.846	0.930	0.930
both arguments are taxa, ≤ 2	33	0.945	0.970	—
both arguments are taxa, ≥ 3	227	0.878	0.947	—
an argument is not a taxon, ≤ 2	128	0.856	0.908	—
an argument is not a taxon, ≥ 3	49	0.628	0.808	—
TEST299	299	0.819	0.922	0.915
BIOTX100	97	0.954	0.970	0.938
REJECT50	41	0.465	0.414	0.503

Table 2: Where the gain lives, AUPRC of the three ensembles. Taxa are counted as distinct mentions found in the passage by TaxoNERD (Le Guillarme and Thuiller, 2022) (en_ner_eco_biobert); an argument “is not a taxon” when TaxoNERD does not recognise it in the passage. TEST299 passages name 5.1 taxa on average, BIOTX100 2.6 and REJECT50 1.7.

is ahead at every one of 99 thresholds, by at least 0.058 F1 (Figure 1).

Adding the relation term to the query does not add to this. The triple-conditioned model is also far ahead of the sentence-level one ($p = 6.8 \times 10^{-5}$), but against the pair-conditioned model it fixes 20 items and breaks 31 ($p = 0.16$), and its AUPRC is, if anything, lower. What the decision needs is the pair. The one place the relation term earns its keep is at strict operating points: at $\tau = 0.5$ both reach precision 0.913, the triple model at recall 0.768 and the pair model at 0.683 (Table 5).

Table 2 locates the gain, and it is the most direct test of the argument. On passages that name three or more taxa, the pair query is worth +0.084 AUPRC. On passages naming two or fewer it is worth less, +0.042, and the split below shows why it is worth anything there at all: the candidates come from automatic matching, and in 41% of

Model	AUPRC	P	F1	AUPRC, sentences naming	
				2 concepts	≥ 3
Sentence-level	0.783 $\pm$.007	0.608	0.725	0.969	0.768
Pair-conditioned	0.876$\pm$.006	0.748	0.834	0.961	0.874

Table 3: BioRED test, 3,277 sentence-level candidates from 100 abstracts, 0.451 positive; 169 come from sentences naming two concepts and 3,108 from sentences naming three or more. AUPRC is the per-seed mean $\pm$ standard deviation over three seeds; P and F1 are at the pre-specified threshold (mean of the checkpoints’ development thresholds). The pair-conditioned model fixes 548 items and breaks 95 (McNemar $p = 4.5 \times 10^{-71}$).

them one “taxon” is a string TaxoNERD does not recognise as an organism in the passage — an organ, an author’s name, a vernacular term. Where both arguments are recognised taxa and the passage names no third one, the case in which the sentence-level question and the pair question coincide, the pair query adds +0.026 on 33 items; where an argument is not a taxon it adds +0.052, and +0.180 when the passage also names others. The pair query does two things, then: it tells the model which of several pairs to judge, and it lets the model check that what it is judging is a pair of organisms. BioRED, whose entities are gold-annotated, isolates the first.

4.2 A public benchmark: BioRED

The biodiversity benchmark is ours, singly annotated, and modest in size. To test the claim where anyone can check it, we rebuild the same setting on BioRED (Luo et al., 2022), a public corpus of 600 PubMed abstracts with expert-annotated entities and relations. Candidates are generated exactly as the biodiversity pipeline generates them, by co-occurrence: every pair of distinct annotated concepts co-mentioned in one sentence, restricted to the eight entity-type pairs that carry relations in the training split, with the sentence as passage. A candidate is positive when BioRED annotates a relation between the two concepts. BioRED’s relations are document-level, so a pair co-mentioned in a sentence that does not itself state the relation still counts as positive; the labels carry that noise, identically for both models. The 400 training abstracts give 11,798 candidates and the 100 test abstracts 3,277. BioRED has no retrieved relation term, so the comparison is between the sentence-level model and the pair-conditioned one, trained with the recipe of §3 on BioRED’s own gold labels.

Query given to the triple model	P	R	F1	AUPRC
Retrieved triple	0.820	0.886	0.852	0.909
Generic relation, pair kept	0.897	0.671	0.767	0.888
Relation shuffled across items	0.881	0.720	0.792	0.904
Relation only, no taxa	0.833	0.020	0.040	0.677
Empty query	1.000	0.037	0.071	0.641

Table 4: Degrading segment A of the trained triple-conditioned model at inference, weights and its own block-held-out thresholds unchanged. Base rate 0.563. Per-seed standard deviations on AUPRC are ≤ 0.027 .

The result (Table 3) reproduces the biodiversity one in both its parts. Conditioning on the pair raises AUPRC from 0.783 to 0.876 and F1 from 0.725 to 0.834, fixing 548 candidates and breaking 95. And the gain sits in the same place: on the 169 candidates from sentences that name only two concepts the two models are indistinguishable (0.969 against 0.961; per seed 0.962–0.965 against 0.954–0.959), while on the 3,108 from sentences naming three or more conditioning is worth +0.106. Ninety-five per cent of BioRED’s co-occurrence candidates come from such sentences, which is why the sentence-level formulation fails there.

4.3 The query, not the encoder

Degrading the trained triple model’s query at inference, with weights and thresholds unchanged, tells the same story from the other side (Table 4). Deleting both taxa collapses ranking from 0.909 to 0.677 against a base rate of 0.563, and emptying the query entirely gives 0.641. Replacing or shuffling the relation, pair kept, barely moves ranking (0.888, 0.904): the model leans on the pair, which is what the trained pair-only model already showed.

The empty-query arm is *not* the sentence-level baseline: it is a model trained with the query and deprived of it, and it does far worse (0.641) than a model trained without the query from the start (0.851). Ablating an input at inference measures dependence on it; it does not stand in for training without it, which is why the comparison of Table 1 trains all three.

Nor does the gain come from the corpus’s pair-contrast design. A triple-conditioned model trained on the 20,000-row base sample alone reaches per-seed AUPRC 0.918 ± 0.003 , against 0.909 ± 0.004 for the full corpus, and the same block-held-out F1 as the pair model, 0.871. The input formulation is what buys the improvement; the extra supervision buys nothing we can measure here.




Figure 1: Precision, recall and F1 of the pair-conditioned model against the decision threshold on all 437 rows, with the sentence-level model’s F1 for comparison (same corpus and encoder). The shaded band spans the three block-held-out thresholds. The pair model is ahead at every τ .

τ	sentence			pair			triple		
	P	R	F1	P	R	F1	P	R	F1
0.01	.690	.915	.787	.745	.976	.845	.721	.935	.814
0.04	.763	.813	.787	.795	.943	.862	.790	.919	.850
0.10	.796	.732	.763	.835	.907	.869	.823	.907	.863
0.25	.837	.687	.754	.889	.850	.869	.877	.841	.859
0.50	.866	.524	.653	.913	.683	.781	.913	.768	.834
0.90	.930	.163	.277	.963	.423	.588	.945	.626	.753

Table 5: The three models’ operating curves on all 437 rows, three-checkpoint ensembles. Rows show the shape; the reported operating points are those of Table 1.

4.4 Recall side

REJECT50 asks whether the candidates the retrieval pipeline discarded contain recoverable positives, and it is too small to rank the models: nine positives, so one item moves recall by eleven points. At block-held-out thresholds the sentence-level model recovers 7 of the 9 by re-admitting 15 of the 32 correct rejections (F1 0.452), the pair-conditioned model recovers 4 while re-admitting 7 (0.400), and the triple-conditioned model recovers 6 while re-admitting 5 (0.600). We read only the qualitative point: the discarded candidates do contain recoverable interactions, and a model that reads the candidate finds them at a lower cost in wrong re-admissions than one that reads only the passage.

4.5 An operating curve, not a verdict

The pair model’s precision moves from 0.745 to 0.963 as τ moves (Table 5), and the right point on

that curve is not a property of the model. Ingestion into a reference database pays for false positives and belongs at the high end; assembling candidates for expert review pays for false negatives and belongs at the low end. A filter that emits a verdict fixes that choice at training time, for users who were never asked. A calibrated score with an adjustable threshold leaves it to them (Geifman and El-Yaniv, 2017).

5 What the remaining errors are made of

The pair-conditioned model’s 39 false positives at its block-held-out operating point are not homogeneous: 27 come from passages naming three or more taxa, and 17 have an argument TaxoNERD does not recognise as a taxon. Read by hand, they sort into a small number of mechanisms (Table 6; the per-row table is released): *pair-binding*, where both taxa and the relation are correct but the relation’s argument is a third organism; *co-listing*, where the two taxa are members of one enumeration and interact with neither each other nor anything in particular; *negation*, where the passage asserts the opposite; *same organism*, where the two “taxa” are one organism named twice or an organism and its own clade; candidates whose relation term names no biotic interaction at all (*relative to*, *hybridization*, *migration*); and *entity failures*, where a taxon string is a taxonomic author, a place, an organ, or an adjective inside a pathogen’s name (“*equine influenza virus*”, which the pair-conditioned model already rejects). A further hand-ful read as gold errors and are discussed in the Limitations.

A bound on component-wise verification. The pair-binding class has a structural property worth stating, because it rules out the obvious alternative architecture. Its members have both taxa correctly resolved *and* the correct relation; what is wrong is a binding that exists only in retrieval’s combinatorics. A verifier that checks s_1 , r and s_2 separately — however good each check is — cannot see it. On the 97 clean BIOTX100 candidates, where triple-level grades make the classes countable, 15 of the 32 unsupported candidates are pair-binding failures (the rest: 9 entity, 8 predicate), a share of 0.47, Clopper–Pearson [0.29, 0.65]. A component-wise verifier accepting a fraction r of the 65 supported triples therefore also accepts all 15, so its precision is at most $65r/(65r + 15)$: 0.81 at full recall, [0.76, 0.88] across that interval. Only a decision

that reads the pair *in context* can go above it.

6 Candidate rules

Several of these mechanisms are not questions about the passage at all but about the candidate, and a classifier is a poor instrument for them: whether two strings name one organism, whether a word is a surname, whether “relative to” is an interaction. We answer those with eight deterministic rejection rules applied to the candidate before the verifier decides. *same_organism* (apposition of a common and a scientific name, or identical forms), *own_clade* (an adjacent clade annotation such as “(Acari: Prostigmata)”), and *authority* (a taxonomic author parsed as a taxon) detect the same-organism and entity classes. *non_biotic_relation* rejects phylogenetic, comparative and biogeographic relation terms; *negated* rejects the two negation forms whose scope provably contains one of the pair (“non- s hosts”, “no infections with ... s ”); *not_an_organism* rejects organs and syndromes; *pathogen_modifier* rejects an adjective that only ever occurs inside a pathogen’s name; and *co_listed* rejects two taxa that are conjuncts of one dependency-parsed list not governed by an interaction noun (“interactions between A and B ” is kept).

How they were validated, and what that licenses. No rule has a parameter fitted to evaluation labels, but rules can still be fitted by *choice*, and four of these — *non_biotic_relation*, *negated*, *not_an_organism* and *pathogen_modifier* — were written after reading the false positives above. To keep their selection from simply memorising the benchmark, every rule had to pass a criterion fixed before it was evaluated: on the 42,747 training rows, at least 85% of the candidates it rejects must be teacher-negative, matching the teacher’s own accuracy. Two first drafts failed it and were narrowed on training data alone — a general negation cue fired on “without affecting the host”, and a role-noun list rejected “viruses” and “fungi”, which are legitimate coarse taxa — and the final rules reject teacher-negatives at 0.88 to 1.00 (*co_listed*, which needs a dependency parse, on a fixed 6,000-row sample). The rules were then frozen, recorded by hash, and evaluated once.

The result (Table 7) is smaller than the motivation suggested, and the reason is the headline of this paper. On the pair-conditioned model the rules buy 0.009 to 0.018 precision for 0.008 to 0.020

Mechanism	Candidate triple	Passage (abridged)
Pair-binding	(<i>T. gondii</i> , parasite of, swine)	“The entero-pathogenic <i>Cystoisospora suis</i> , a coccidian parasite of swine and close relative to <i>Toxoplasma gondii</i> ...”
Co-listing	(katydid, prey, beetles)	“... remains were from insects ... (e.g. beetles , dragonflies, cockroaches, and female katydids), indicating that eavesdropping on prey signals is not the only prey-finding strategy used by this <i>bat</i> .”
Negation	(wheat, hosts, <i>F. acuminatum</i>)	“... <i>F. acuminatum</i> in New Zealand was found only on non-wheat hosts .”
Same organism	(Chinese knotweed, pollinators, <i>P. chinensis</i>)	“... potential pollinators of Chinese knotweed (<i>Persicaria chinensis</i>).”
Entity	(Siberia , migrating, birds)	“... a portion of birds migrating from Siberia have common winter nesting areas.”

Table 6: One example of each mechanism among the pair-conditioned model’s false positives. The relation’s true argument is italicised where it is not the queried taxon; in the entity case a place name was matched as a taxon.

System	P	R	F1	fixed	broken
Pair-conditioned	0.849	0.894	0.871	—	—
+ 7 rules	0.858	0.886	0.872	3	2
+ 8 rules	0.867	0.874	0.870	6	5
Triple-conditioned	0.820	0.886	0.852	—	—
+ 7 rules	0.840	0.878	0.859	7	2
+ 8 rules	0.862	0.866	0.864	14	5

Table 7: Deterministic candidate rules applied before the model’s decision, all 437 rows, block-held-out thresholds. “Fixed” and “broken” are relative to the same model without rules. McNemar: $p = 1.00$ and 1.00 on the pair model, 0.18 and 0.067 on the triple model.

recall, F1 unchanged, and none of it is distinguishable from no effect: of its 39 false positives the eight rules catch 6. A model told which pair it is judging already rejects most of what the rules were written to reject — a non-organism argument, a co-listed pair — because it can read the candidate. On the triple-conditioned model, which leans on its relation term, the rules do more: precision $0.820 \rightarrow 0.862$ with 14 items fixed and 5 broken ($p = 0.067$). We report these as upper estimates in any case, since half the rules were chosen with this benchmark’s errors in view; what the training-data criterion establishes is that they reject what they are meant to reject on 42k rows they were not chosen on, not that their effect here would replicate exactly. Their case in deployment is transparency rather than accuracy: each rejection names the rule that made it, and a curator can audit a rule where they cannot audit a weight.

7 The deployed configuration

In the pipeline, verification runs as one 110M-parameter model with two heads. It differs from the pair-conditioned model of §3 in three respects,

Expert direction label	n	answered	correct
forward or reverse	84	62 (74%)	55 of 62 (0.887)
undecidable (“?”)	12	—	9 left UNCERTAIN
bidirectional	1	—	1 BIDIRECTIONAL

Table 8: The deployed model’s four-way direction output against 97 expert direction labels. “Answered” counts FORWARD or REVERSE; the rest are UNCERTAIN. The abstention level (0.60) was chosen by the expert after seeing this set’s coverage curve, so the figures are optimistic.

each of which we adopted during development: the two taxa are given in alphabetical order, so the input is identical whichever way round retrieval presents the pair; they are marked in place in the passage (@...@, #...#) rather than only named in segment A; and the training labels are the teacher’s pair-level judgments. Its interaction threshold, 0.5, was set before evaluation and never tuned. On the 437 rows it reaches AUPRC 0.934, and precision 0.851, recall 0.955, F1 0.900; with the candidate rules, precision 0.869, recall 0.947, F1 0.907 (6 items fixed, 2 broken, $p = 0.29$).¹ It needs no GPU: on eight CPU threads it scores 32.7 candidates per second with both heads, after a one-off 7.5 s to load, which is faster than the pipeline retrieves them, so every candidate is verified rather than a sample.

The second head reports the interaction’s *direction* in four categories. FORWARD and REVERSE say which taxon acts on the other; they come from a direction score that is exactly antisymmetric in the two marked taxa, and are reported only when it is confident, otherwise the answer is UNCERTAIN.

¹The *co-listed* rule is omitted from this configuration: on the deployed model’s accept set it removes one false positive at the cost of three true ones.

BIDIRECTIONAL is assigned when the relation is mutual by definition — mutualism, symbiosis, competition, co-infection, *interacts with* — as listed in a relation-polarity lexicon derived from the Relation Ontology; for those, neither taxon is the subject and the model is not asked. Expert curation of the benchmark scores reversed triples as correct (§8), so direction is evaluated on a separate expert annotation of 97 candidates (Table 8): the model answers 74% of the 84 with a direction and is right on 0.887 of those, leaves 9 of the 12 the expert could not decide as UNCERTAIN, and labels the one bidirectional item BIDIRECTIONAL.

8 Where the reformulation stops

Entity disambiguation, and a gate that cannot fix it. The natural remedy for entity errors is to gate candidates on taxonomic validity. Against the Open Tree Taxonomy (Rees and Cranston, 2017) such a gate catches 1 of the 35 unsupported BIOTX100 candidates and discards 5 of the 65 supported ones. The reason is diagnostic: *Illiger*, *Dikarya*, *copia* and *sulla* — the surnames, clades, transposon families and vernacular names behind the entity failures — all *resolve* in the taxonomy. The failure is disambiguation, of the kind species-name normalisers address (Gerner et al., 2010), not validity; and the gate’s casualties are viruses, which are legitimate partners in this domain. The *authority* and *same_organism* rules (§6) target the disambiguation directly and are the better instrument.

Direction needs its own gold. Expert curation of the benchmark marks reversed triples as correct — we verified two by hand, (*Breinlia*, host of, *Rattus*) and (*Demidospermus*, parasitized by, *A. inermis*) — so no directional number on the benchmark would measure the model rather than the convention, which is why direction is scored on its own annotation (§7). The supervision itself is uneven in a revealing way: among the 5,591 reversed training rows, the rate at which the teacher rejects a reversal is governed by the relation’s *surface string*. Direction-encoding participles and prepositional forms are rejected at 0.51–0.70 (*affecting* 0.70, *affected by* 0.61, *parasite of* 0.57), bare relational nouns at 0.04–0.10 (*host of* 0.08, *pollinators* 0.04). Whether supervision encodes order at all depends on which string retrieval happened to match.

9 Related work

Classifiers that answer an easier question. That a classifier can score well by exploiting how its data were assembled is among the best-established results in NLP. Partial-input baselines recover much of the label from the hypothesis alone in natural language inference (Gururangan et al., 2018; Poliak et al., 2018) and from the passage alone in reading comprehension (Kaushik and Lipton, 2018), and such rules hold in distribution while failing under deployment (Geirhos et al., 2020; Du et al., 2024); the closest task-level precedent is fact verification (Thorne et al., 2018), where a claim-only classifier matches evidence-aware models because of how the corpus was built (Schuster et al., 2019). Our sentence-level baseline is a partial-input model in exactly this sense, and the specific feature here is that the missing input is the query the pipeline itself issued rather than an annotator artefact (Wiegand et al., 2019; Parmar et al., 2023). Feng et al. (2019) caution that partial-input results license claims about models and corpora rather than about tasks, and we keep ours to the former.

The argument-ignoring failure is already named. Distant supervision assumes that a sentence containing two related entities expresses the relation (Mintz et al., 2009), and its standard relaxation exists because it does not (Riedel et al., 2010). Since TACRED an instance has been the triple (sentence, e_1 , e_2) rather than the sentence (Zhang et al., 2017), and how the pair is injected into the encoder is worth several F1 points (Baldini Soares et al., 2019; Wu and He, 2019; Zhou and Chen, 2022). Rosenman et al. (2020) name our failure exactly, as an “event heuristic” that classifies on (sentence, relation) while ignoring the arguments. *We claim no part of that diagnosis.* What is ours is a controlled measurement of what it costs in a verification setting built by co-occurrence retrieval, located by block; a bound on component-wise remedies; and the candidate rules. The entity-type shortcut usually invoked for such errors (Peng et al., 2020; Mtumbuka and Schockaert, 2024) is unavailable here, since both slots are filled from one inventory of taxon surface forms. Error attribution across pipeline stages is standard and easy to get wrong (Taillé et al., 2020; Ai and Kavuluru, 2023).

Conditioning on the query. Casting a relation as a question about a named entity (Levy et al., 2017), or as an entailment hypothesis verbalising

the relation between two of them (Obamuyide and Vlachos, 2018; Sainz et al., 2021), is the reformulation our input format instantiates, and the resulting decision — does this passage support this specific claim — has the shape of scientific claim verification (Wadden et al., 2020), and of validating a knowledge-graph triple against text (Lehmann et al., 2012), with the difference that those systems retrieve evidence for a given triple while here the passage is fixed by the retrieval that proposed the triple and the question is whether that passage, and no other, supports it. Distilling a large model’s judgments into a small student is established practice (West et al., 2022; Gekhman et al., 2023). Selective classification thresholds a score instead of issuing a verdict (Geifman and El-Yaniv, 2017), which is how extraction is made useful to human abstractors: Swaminathan et al. (2024) abstain under calibrated thresholds to automate 57–95% of clinical chart abstraction at controlled error.

Text mining for interaction databases. GloBI (Poelen et al., 2014) is populated in part by text mining (Dimitrova et al., 2020). Judging an extracted fact before it reaches a curator is an old idea (Rodriguez-Esteban et al., 2006), as is evaluating such tools by whether they help the curator (Hirschman et al., 2012; Britan et al., 2018; Lee et al., 2018). Cuzick et al. (2023), independently, take the interacting pair rather than the passage as the annotation unit — the unit our pair-binding class isolates. Large language models are applied to interaction extraction at scale (Keck et al., 2025; Farrell et al., 2024; D’Souza et al., 2025; Zou et al., 2026), but as extractors: Zou et al. (2026) standardise names and categories post hoc without checking extracted interactions against the text that produced them. We are not aware of published work that verifies literature-mined biotic-interaction triples against their source passage before ingestion.

10 Conclusion

When candidates are retrieved by co-occurrence, the sentence-level question is answered before it is asked, and a classifier built for it is left with little to decide. Telling the model which pair the candidate is about — with the same encoder and the same data — raises AUPRC from 0.851 to 0.919 on the biodiversity benchmark and from 0.783 to 0.876 on BioRED, and the relation term adds nothing on top. The gain sits where the argument puts it: on passages that name more than one candidate

pair, and, where candidates come from automatic matching, on candidates whose arguments are not organisms. The result is a calibrated score where a verdict stood, with a four-way direction label beside it, cheap enough to run on every candidate.

Limitations

One benchmark, singly annotated. 437 rows in three blocks, every label assigned by one domain expert, with no independent second grading: we report no inter-annotator agreement and cannot bound label noise. Some apparent false positives are not model errors in the usual sense. A passage stating that “the bat tick . . . has been reported . . . to occasionally feed on humans” is labelled negative, correctly, because the curator judged the candidate’s *resolved* taxon, which retrieval had mapped to the wrong species; the model sees only the surface form, which the passage does support. Errors of this kind are upstream, and no passage-reading model can repair them. Candidates on which two independent models confidently contradict the gold go to blind expert re-annotation, mixed with an equal number of random controls and with labels and scores hidden; no label is changed on the strength of a model’s disagreement. On the recall side the denominator is 9 positives.

Top-ranked candidates only. Every BIOTX100 row is the top-ranked candidate for its passage, and half are pairs GloBI already records. Whether the gains hold further down the retrieval ranking, where candidates are noisier, is not measured here; TEST299 and the public benchmark of §4.2 are not rank-restricted.

The rules are partly fitted by choice. Four of the eight rules were selected after reading this benchmark’s false positives (§6). The training-data criterion bounds how much that selection can have memorised the benchmark, but it does not remove it, and their effects should be read as optimistic until confirmed on candidates neither the rules nor their author have seen. They fire on no row of a separate 100-candidate sample, so that sample can neither confirm nor contradict them.

Fresh candidates are harder. On that sample — 100 retrieved candidates drawn after the benchmark, half of them from passages containing several candidate pairs — all three models score well below their benchmark figures, though in the same

order: AUPRC 0.665 sentence-level, 0.785 pair-conditioned, 0.765 triple-conditioned. Its label provenance is not identical to the benchmark’s, so we do not read the drop as measured, but it is the clearest sign that multi-pair passages remain the hard case and that the benchmark may understate their share of a live pool.

The diagnosis is triple-level; the evaluation is pair-level. Only BIOTX100 carries whole-triple gold. Evaluating at the pair level everywhere means a candidate with the right pair and a misbound relation counts as correct, so the comparison cannot credit the relation term with the cases it alone could catch; triple-level gold on the other two blocks is what would test whether the triple model’s advantage at strict operating points is real.

BioRED’s labels are document-level. A BioRED candidate is positive when the abstract relates the two concepts anywhere, so a co-mention in a sentence that does not state the relation counts as positive. Both models are trained and scored on the same noisy labels, so the comparison is controlled, but the absolute figures understate what either would reach against sentence-level gold.

Teacher noise and pretraining exposure. The training labels come from a prompted teacher wrong about one time in seven, and the student inherits that boundary; training on model-generated labels can degrade encoders (Lu and Smith, 2025). The passages are published MEDLINE text and the teacher is a pretrained model, so we cannot exclude that it saw them, and its labels may be optimistic for genuinely new literature.

Upstream failures are out of reach. Of the 50 candidates the retrieval pipeline discarded, 32 carry an incorrect interaction term and 17 a mis-resolved taxon, 14 of them both. Those fail before any verifier is consulted, and no change to the decision stage, ours included, addresses them.

One system, one domain. One retrieval system, one interaction vocabulary, taxon–taxon interactions in MEDLINE. The mechanism — a pair-agnostic decision over passages that name several candidate pairs — is not specific to ours, but we have not shown that other pools fail the same way.

References

- Xuguang Ai and Ramakanth Kavuluru. 2023. End-to-end models for chemical–protein interaction extraction: Better tokenization and span-based pipeline strategies. In *2023 IEEE 11th International Conference on Healthcare Informatics (ICHI)*, pages 610–618. IEEE. ArXiv:2304.01344.
- Livio Baldini Soares, Nicholas FitzGerald, Jeffrey Ling, and Tom Kwiatkowski. 2019. [Matching the blanks: Distributional similarity for relation learning](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2895–2905, Florence, Italy. Association for Computational Linguistics.
- Aurore Britan, Isabelle Cusin, Valérie Hinard, Luc Motin, Emilie Pasche, Julien Gobeill, Valentine Rech de Laval, Anne Gleizes, Daniel Teixeira, Pierre-André Michel, Patrick Ruch, and Pascale Gaudet. 2018. [Accelerating annotation of articles via automated approaches: evaluation of the neXtA5 curation-support tool by neXtProt](#). *Database*, 2018:bay129.
- Cristian Bucilua, Rich Caruana, and Alexandru Niculescu-Mizil. 2006. Model compression. In *Proceedings of the 12th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, pages 535–541.
- Alayne Cuzick, James Seager, Valerie Wood, Martin Urban, Kim Rutherford, and Kim E. Hammond-Kosack. 2023. [A framework for community curation of interspecies interactions literature](#). *eLife*, 12:e84658.
- Mariya Dimitrova, Jorrit Poelen, Georgi Zhelezov, Teodor Georgiev, Donat Agosti, and Lyubomir Penev. 2020. [Semantic publishing enables text mining of biotic interactions](#). *Biodiversity Information Science and Standards*, 4:e59036.
- Jennifer D’Souza, Zachary Laubach, Tarek Al Mustafa, Sina Zarriß, Robert Frühstückl, and Phyllis Illari. 2025. [Mining for species, locations, habitats, and ecosystems from scientific papers in invasion biology: A large-scale exploratory study with large language models](#). In *Proceedings of the 1st Workshop on Ecology, Environment, and Natural Language Processing (NLP4Ecology2025)*, pages 16–23, Tallinn, Estonia. University of Tartu Library.
- Mengnan Du, Fengxiang He, Na Zou, Dacheng Tao, and Xia Hu. 2024. [Shortcut learning of large language models in natural language understanding](#). *Communications of the ACM*, 67(1):110–120.
- Maxwell J. Farrell, Nicolas Le Guillarme, Liam Brierley, Bronwen Hunter, Daan Scheepens, Anna Willoughby, Andrew Yates, and Nicole Mideo. 2024. [The changing landscape of text mining: a review of approaches for ecology and evolution](#). *Proceedings of the Royal Society B: Biological Sciences*, 291(2027):20240423.
- Shi Feng, Eric Wallace, and Jordan Boyd-Graber. 2019. [Misleading failures of partial-input baselines](#). In

841	<i>Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics</i> , pages 5533–5538, Florence, Italy. Association for Computational Linguistics.	
842		
843		
844		
845	Yonatan Geifman and Ran El-Yaniv. 2017. Selective classification for deep neural networks. In <i>Advances in Neural Information Processing Systems 30 (NIPS 2017)</i> , pages 4878–4887.	
846		
847		
848		
849	Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. 2020. Shortcut learning in deep neural networks . <i>Nature Machine Intelligence</i> , 2(11):665–673.	
850		
851		
852		
853		
854	Zorik Gekhman, Jonathan Herzig, Roei Aharoni, Chen Elkind, and Idan Szpektor. 2023. TrueTeacher: Learning factual consistency evaluation with large language models . In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing</i> , pages 2053–2070, Singapore. Association for Computational Linguistics.	
855		
856		
857		
858		
859		
860		
861	Martin Gerner, Goran Nenadic, and Casey M. Bergman. 2010. LINNAEUS: A species name identification system for biomedical literature . <i>BMC Bioinformatics</i> , 11:85.	
862		
863		
864		
865	Julien Gobeill, Déborah Caucheteur, Pierre-André Michel, Luc Mottin, Emilie Pasche, and Patrick Ruch. 2020. SIB literature services: RESTful customizable search engines in biomedical literature, enriched with automatically mapped biomedical concepts . <i>Nucleic Acids Research</i> , 48(W1):W12–W16.	
866		
867		
868		
869		
870		
871	Yu Gu, Robert Tinn, Hao Cheng, Michael Lucas, Naoto Usuyama, Xiaodong Liu, Tristan Naumann, Jianfeng Gao, and Hoifung Poon. 2021. Domain-specific language model pretraining for biomedical natural language processing. <i>ACM Transactions on Computing for Healthcare</i> , 3(1):1–23.	
872		
873		
874		
875		
876		
877	Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. 2018. Annotation artifacts in natural language inference data . In <i>Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)</i> , pages 107–112, New Orleans, Louisiana. Association for Computational Linguistics.	
878		
879		
880		
881		
882		
883		
884		
885		
886	Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015. Distilling the knowledge in a neural network. In <i>NIPS Deep Learning and Representation Learning Workshop</i> .	
887		
888		
889		
890	Lynette Hirschman, Gully A. P. C. Burns, Martin Krallinger, Cecilia Arighi, K. Bretonnel Cohen, Alfonso Valencia, Cathy H. Wu, Andrew Chatr-Aryamontri, Karen G. Dowell, Eva Huala, Anália Lourenço, Robert Nash, Anne-Lise Veuthey, Thomas Wiegers, and Andrew G. Winter. 2012. Text mining for the biocuration workflow . <i>Database</i> , 2012:bas020.	
891		
892		
893		
894		
895		
896		
897		
	Divyansh Kaushik and Zachary C. Lipton. 2018. How much reading does reading comprehension require? A critical investigation of popular benchmarks . In <i>Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing</i> , pages 5010–5015, Brussels, Belgium. Association for Computational Linguistics.	898
		899
		900
		901
		902
		903
		904
	François Keck, Henry Broadbent, and Florian Altermatt. 2025. Extracting massive ecological data on state and interactions of species using large language models . <i>bioRxiv</i> . Preprint, posted 27 January 2025.	905
		906
		907
		908
	Nicolas Le Guillarme and Wilfried Thuiller. 2022. TaxoNERD: deep neural models for the recognition of taxonomic entities in the ecological and evolutionary literature . <i>Methods in Ecology and Evolution</i> , 13(3):625–641.	909
		910
		911
		912
		913
	Kyubum Lee, Maria Livia Famiglietti, Aoife McMahon, Chih-Hsuan Wei, Jacqueline Ann Langdon MacArthur, Sylvain Poux, Lionel Breuza, Alan Bridge, Fiona Cunningham, Ioannis Xenarios, and Zhiyong Lu. 2018. Scaling up data curation using deep learning: An application to literature triage in genomic variation resources . <i>PLOS Computational Biology</i> , 14(8):e1006390.	914
		915
		916
		917
		918
		919
		920
		921
	Jens Lehmann, Daniel Gerber, Mohamed Morsey, and Axel-Cyrille Ngonga Ngomo. 2012. DeFacto – deep fact validation. In <i>The Semantic Web – ISWC 2012</i> , volume 7649 of <i>Lecture Notes in Computer Science</i> , pages 312–327. Springer.	922
		923
		924
		925
		926
	Omer Levy, Minjoon Seo, Eunsol Choi, and Luke Zettlemoyer. 2017. Zero-shot relation extraction via reading comprehension . In <i>Proceedings of the 21st Conference on Computational Natural Language Learning (CoNLL 2017)</i> , pages 333–342, Vancouver, Canada. Association for Computational Linguistics.	927
		928
		929
		930
		931
		932
	Yucheng Lu and Kazimier Smith. 2025. Feeding LLM annotations to BERT classifiers at your own risk. <i>arXiv preprint arXiv:2504.15432</i> .	933
		934
		935
	Ling Luo, Po-Ting Lai, Chih-Hsuan Wei, Cecilia N Arighi, and Zhiyong Lu. 2022. BioRED: a rich biomedical relation extraction dataset. <i>Briefings in Bioinformatics</i> , 23(5):bbac282.	936
		937
		938
		939
	Quinn McNemar. 1947. Note on the sampling error of the difference between correlated proportions or percentages. <i>Psychometrika</i> , 12(2):153–157.	940
		941
		942
	Mike Mintz, Steven Bills, Rion Snow, and Dan Jurafsky. 2009. Distant supervision for relation extraction without labeled data. In <i>Proceedings of ACL-IJCNLP</i> , pages 1003–1011.	943
		944
		945
		946
	Frank Martin Mtumbuka and Steven Schockaert. 2024. Entity or relation embeddings? an analysis of encoding strategies for relation extraction . In <i>Findings of the Association for Computational Linguistics: EMNLP 2024</i> , pages 6003–6022, Miami, Florida, USA. Association for Computational Linguistics.	947
		948
		949
		950
		951
		952

953	Rodrigo Nogueira and Kyunghyun Cho. 2019. Pas-	1009
954	sage re-ranking with BERT . <i>arXiv preprint</i>	1010
955	<i>arXiv:1901.04085</i> .	1011
956	Abiola Obamuyide and Andreas Vlachos. 2018. Zero-	1012
957	shot relation classification as textual entailment . In	
958	<i>Proceedings of the First Workshop on Fact Extraction</i>	1013
959	<i>and VERification (FEVER)</i> , pages 72–78, Brussels,	1014
960	Belgium. Association for Computational Linguistics.	1015
961	Mihir Parmar, Swaroop Mishra, Mor Geva, and Chitta	1016
962	Baral. 2023. Don’t blame the annotator: Bias al-	1017
963	ready starts in the annotation instructions . In <i>Pro-</i>	1018
964	<i>ceedings of the 17th Conference of the European</i>	1019
965	<i>Chapter of the Association for Computational Lin-</i>	1020
966	<i>guistics</i> , pages 1779–1789, Dubrovnik, Croatia. As-	
967	sociation for Computational Linguistics.	
968	Hao Peng, Tianyu Gao, Xu Han, Yankai Lin, Peng	1021
969	Li, Zhiyuan Liu, Maosong Sun, and Jie Zhou. 2020.	1022
970	Learning from context or names? an empirical study	1023
971	on neural relation extraction . In <i>Proceedings of the</i>	1024
972	<i>2020 Conference on Empirical Methods in Natural</i>	1025
973	<i>Language Processing (EMNLP)</i> , pages 3661–3672,	1026
974	Online. Association for Computational Linguistics.	1027
975	Jorrit H Poelen, James D Simons, and Chris J Mungall.	1028
976	2014. Global biotic interactions: An open infrastruc-	1029
977	ture to share and analyze species-interaction datasets.	
978	<i>Ecological Informatics</i> , 24:148–159.	
979	Adam Poliak, Jason Naradowsky, Aparajita Haldar,	1030
980	Rachel Rudinger, and Benjamin Van Durme. 2018.	1031
981	Hypothesis only baselines in natural language infer-	1032
982	ence . In <i>Proceedings of the Seventh Joint Confer-</i>	1033
983	<i>ence on Lexical and Computational Semantics</i> , pages	1034
984	180–191, New Orleans, Louisiana. Association for	1035
985	Computational Linguistics.	1036
986	Qwen Team. 2025. Qwen3 technical report. <i>arXiv</i>	1037
987	<i>preprint arXiv:2505.09388</i> .	1038
988	Jonathan A. Rees and Karen Cranston. 2017. Auto-	1039
989	mated assembly of a reference taxonomy for phylo-	
990	genetic data synthesis . <i>Biodiversity Data Journal</i> ,	1040
991	5:e12581.	1041
992	Sebastian Riedel, Limin Yao, and Andrew McCallum.	1042
993	2010. Modeling relations and their mentions with-	1043
994	out labeled text . In <i>Machine Learning and Knowl-</i>	1044
995	<i>edge Discovery in Databases (ECML PKDD 2010)</i> ,	1045
996	Lecture Notes in Computer Science, pages 148–163,	1046
997	Barcelona, Spain. Springer Berlin Heidelberg.	
998	Raul Rodriguez-Esteban, Ivan Iossifov, and Andrey	1047
999	Rzhetsky. 2006. Imitating manual curation of text-	1048
1000	mined facts in biomedicine . <i>PLoS Computational</i>	1049
1001	<i>Biology</i> , 2(9):e118.	1050
1002	Shachar Rosenman, Alon Jacovi, and Yoav Goldberg.	1051
1003	2020. Exposing shallow heuristics of relation ex-	
1004	traction models with challenge data . In <i>Proceed-</i>	1052
1005	<i>ings of the 2020 Conference on Empirical Methods</i>	1053
1006	<i>in Natural Language Processing (EMNLP)</i> , pages	1054
1007	3702–3710, Online. Association for Computational	1055
1008	Linguistics.	1056
	Patrick Ruch, Emilie Pasche, Pierre-André Michel, and	1057
	Déborah Caucheteur. 2024. BiotXplorer: Navigat-	1058
	ing evidence-based biotic interactions . <i>Biodiversity</i>	
	<i>Information Science and Standards</i> , 8:e135453.	
	Oscar Sainz, Oier Lopez de Lacalle, Gorka Labaka,	1059
	Ander Barrena, and Eneko Agirre. 2021. Label ver-	1060
	balization and entailment for effective zero and few-	1061
	shot relation extraction . In <i>Proceedings of the 2021</i>	1062
	<i>Conference on Empirical Methods in Natural Lan-</i>	1063
	<i>guage Processing</i> , pages 1199–1212, Online and	1064
	Punta Cana, Dominican Republic. Association for	
	Computational Linguistics.	
	Tal Schuster, Darsh Shah, Yun Jie Serene Yeo, Daniel	1065
	Roberto Filizzola Ortiz, Enrico Santus, and Regina	1066
	Barzilay. 2019. Towards debiasing fact verification	1067
	models . In <i>Proceedings of the 2019 Conference on</i>	1068
	<i>Empirical Methods in Natural Language Processing</i>	1069
	<i>and the 9th International Joint Conference on Natu-</i>	1070
	<i>ral Language Processing (EMNLP-IJCNLP)</i> , pages	1071
	3419–3425, Hong Kong, China. Association for Com-	1072
	putational Linguistics.	1073
	Akshay Swaminathan, Ivan Lopez, William Wang,	1074
	Ujwal Srivastava, Edward Tran, Aarohi Bhargava-	1075
	Shah, Janet Y. Wu, Alexander L. Ren, Kaitlin Caoili,	1076
	Brandon Bui, Layth Alkhani, Susan Lee, Nathan	1077
	Mohit, Noel Seo, Nicholas Macedo, Winson Cheng,	1078
	Charles Liu, Reena Thomas, Jonathan H. Chen,	1079
	and Olivier Gevaert. 2024. Selective prediction	1080
	for extracting unstructured clinical data . <i>Journal</i>	1081
	<i>of the American Medical Informatics Association</i> ,	1082
	31(1):188–197.	1083
	Bruno Taillé, Vincent Guigue, Geoffrey Scouteeten,	1084
	and Patrick Gallinari. 2020. Let’s stop incorrect	1085
	comparisons in end-to-end relation extraction! In	1086
	<i>Proceedings of the 2020 Conference on Empirical</i>	1087
	<i>Methods in Natural Language Processing (EMNLP)</i> ,	1088
	pages 3689–3701. Association for Computational	1089
	Linguistics.	1090
	James Thorne, Andreas Vlachos, Christos	1091
	Christodoulopoulos, and Arpit Mittal. 2018.	1092
	FEVER: a large-scale dataset for fact extraction and	1093
	VERification. In <i>Proceedings of NAACL-HLT</i> , pages	1094
	809–819.	1095
	David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu	1096
	Wang, Madeleine van Zuylen, Arman Cohan, and	1097
	Hannaneh Hajishirzi. 2020. Fact or fiction: Verifying	1098
	scientific claims . In <i>Proceedings of the 2020 Con-</i>	1099
	<i>ference on Empirical Methods in Natural Language</i>	1100
	<i>Processing (EMNLP)</i> , pages 7534–7550, Online. As-	1101
	sociation for Computational Linguistics.	1102
	Peter West, Chandra Bhagavatula, Jack Hessel, Jena D.	1103
	Hwang, Liwei Jiang, Ronan Le Bras, Ximing Lu,	1104
	Sean Welleck, and Yejin Choi. 2022. Symbolic	1105
	knowledge distillation: from general language mod-	1106
	els to commonsense models . In <i>Proceedings of the</i>	1107
	<i>2022 Conference of the North American Chapter of</i>	1108

the Association for Computational Linguistics: Human Language Technologies, pages 4602–4625, Seattle, United States. Association for Computational Linguistics.

Michael Wiegand, Josef Ruppenhofer, and Thomas Kleinbauer. 2019. [Detection of abusive language: the problem of biased datasets](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 602–608, Minneapolis, Minnesota. Association for Computational Linguistics.

Shanchan Wu and Yifan He. 2019. [Enriching pre-trained language model with entity information for relation classification](#). In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management (CIKM)*, pages 2361–2364, Beijing, China. Association for Computing Machinery.

Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli, and Christopher D. Manning. 2017. [Position-aware attention and supervised data improve slot filling](#). In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 35–45, Copenhagen, Denmark. Association for Computational Linguistics.

Wenxuan Zhou and Muhao Chen. 2022. [An improved baseline for sentence-level relation extraction](#). In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 161–168, Online only. Association for Computational Linguistics.

Heng-Xing Zou, Xiaohao Yang, Thabassum H. Hajamaideen, Olivia J. Stein, Roxanne S. Beltran, Benjamin G. Freeman, Mark Lindquist, Eliot T. Miller, Summer Mengarelli, Charlotte M. Probst, Fernanda S. Valdovinos, Derek Van Berkel, Phoebe L. Zarnetske, Brian C. Weeks, and Kai Zhu. 2026. [Large language models unlock the ecology of species interactions](#). *bioRxiv*. Preprint, posted 10 February 2026.

A Per-seed results

B Negative results and cost

Table 10 collects interventions that did not pay, because each is a plausible first thing to try. Self-distillation at equal capacity is null, suggesting the teacher’s knowledge is already saturated in the student. Compression to six layers buys throughput the pipeline does not need for 0.088 precision. Reweighting the corpus toward the deployment signal distribution cost 0.059 AUPRC. Most instructive is the last row: before the candidate rules existed, we turned the same mechanisms into *training data* — synthetic negatives for the same-organism,

seed	AUPRC			F1, $\tau = 0.5$		
	sent.	pair	triple	sent.	pair	triple
1	.858	.922	.912	.632	.790	.842
2	.851	.921	.905	.698	.758	.847
3	.845	.912	.909	.663	.788	.834
mean	.851	.919	.909	.664	.779	.841
sd	.006	.005	.004	.033	.018	.006

Table 9: Per-seed figures on the 437 biodiversity rows, order-free scoring. The body reports three-checkpoint ensembles at block-held-out thresholds; a fixed $\tau = 0.5$ is shown here because the three models are calibrated differently, and a single shared threshold would favour one of them arbitrarily.

Intervention	Benefit	Cost
Taxon-validity gate	1 of 35 unsupported caught	−5/65 supported; rejects viruses
Self-distillation	none	−0.001 P, −0.008 F1
6-layer compression	1.9× CPU throughput	−0.088 precision
Corpus signal-matching	none	−0.059 AUPRC
Rule-family training data	none	AUPRC +0.005 to −0.037

Table 10: Five plausible design choices that measurably do not pay.

own-clade and authority classes, and in a second variant for co-listing — and retrained. The first variant left AUPRC unchanged within seed noise and the second cost 0.037. Mechanisms that are properties of the candidate are better answered by a rule than taught to a classifier. Deltas in the first four rows were measured under an earlier version of the evaluation and are given as directions.